



University
of Glasgow

Introduction to Generative Adversarial Networks.

Adalberto Claudio Quiros - a.claudio-quiros.1@research.gla.ac.uk

April 22, 2020

Computing Science Department

What's a GAN? Where do they fall in the Machine Learning perspective?

- Deep Learning Model.
- Unsupervised Learning: Generative Models.

What's the goal of a GAN then?

- Learn a distribution of data whether if it is images, sound, vectors, anything.
- Learn it in an unsupervised way: no label or initial knowledge about the data, just samples.

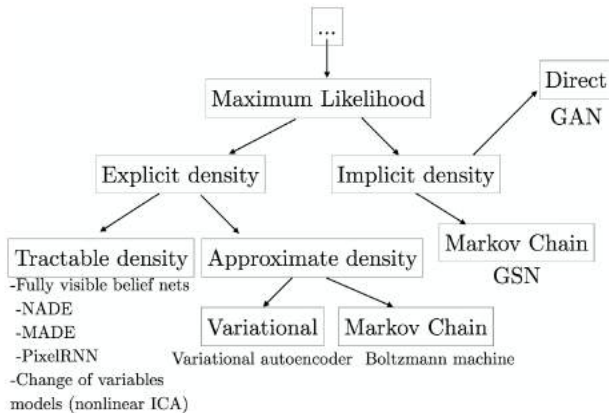



Figure 1: Generative Models - Goodfellow NIPS 2016 Tutorial: Generative Adversarial Networks.

Model Diagram:



Core Idea:

- Two modules: Generator G and Discriminator D .
- The Generator G and Discriminator D compete between each other. G to recreate real data samples and D to distinguish between real and generated samples.

1. Generator:

- **Goal:** Replicate real data.
- **How:** Modifies an prior distribution $p_z(z)$ (latent space) into a distribution $p_g(x)$, making it as close as possible to the real data distribution $p_{data}(x)$
- $G(z; \theta_g)$: Function that maps a latent sample $z = \{z_1, \dots, z_m\} \sim p_z(z)$ to a data sample $x = \{x_1, \dots, x_n\} \sim p_g(x)$.

$$G : z \rightarrow x$$

$$\text{where } z \in \mathfrak{R}^m; x \in \mathfrak{R}^n$$

2. Discriminator:

- **Goal:** Distinguish between real and fake samples.
- **How:** Measures the probability of a sample x_{real} or x_{fake} of coming from the real data distribution $p_{data}(x)$.
- $D(x; \theta_d)$: Function that maps a data sample x to a scalar d . $d = 0$ if the discriminator assumes $x = x_{fake}$ and $d = 1$ if the discriminator assumes $x = x_{real}$.

$$D : x \rightarrow d$$

$$\text{where } x \in \mathfrak{R}^n; d \in \mathfrak{R}$$

3. Generator G and Discriminator D compete to 'trick' each other. Discriminator pushes Generator to replicate the real data distribution $p_{data}(x)$

Loss Function:

- Cross entropy loss.
- Generator G tries to minimize accuracy of the Discriminator D .

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Discriminator D :

- First part: Improve the probability of recognizing real data.
- Second part: Improve the probability of recognizing fake data.

$$\max_D V(D) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$


Generator G :

- Improve quality of the generated samples so they will fool the Discriminator D .

$$\min_G V(G) = \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Discriminator D and Generator G gradients:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m [\log D(x^i) + \log(1 - D(G(z^i)))] \text{ and } -\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m [\log(1 - D(G(z^i)))]$$



Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log \left(1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left(1 - D(G(\mathbf{z}^{(i)})) \right).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Figure 2: Training Algorithm Goodfellow et al.[1].

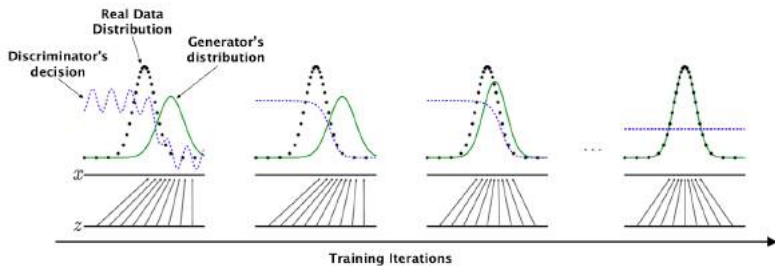


Figure 3: Generated data distribution over training [1].

Let's see it: GAN training

Let's assume a fixed Generator $G(z)$, and find the optimal Discriminator $D^*(x)$, the best Discriminator that we can train for a fixed Generator (**This is what we do in the first part of the algorithm !**):

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

If Discriminator $D(x)$ wants to maximize the loss function $V(D, G)$, which Discriminator achieves this maximum:

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] =$$

$$\int_x p_{data}(x) \log D(x) dx + \int_z p_z(z) \log(1 - D(G(z))) dz =$$

$$\int_x (p_{data}(x) \log D(x) + p_g(x) \log[1 - D(x)]) dx$$

$$\frac{\partial V(D, G)}{\partial D} = 0 \rightarrow p_{data}(x) \frac{1}{D(x)} - p_g(x) \frac{1}{1 - D(x)} = 0$$

$$D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$$

What's the loss function for the Optimal Discriminator $D^*(x)$?

$$\begin{aligned}
 \min_G V(D^*, G) &= \int_x (p_{data}(x) \log D^*(x) + p_g(x) \log(1 - D^*(x))) dx = \\
 &\int_x (p_{data}(x) \log \left[\frac{p_{data}(x)}{p_{data}(x) + p_g(x)} \right] + p_g(x) \log \left[\frac{p_g(x)}{p_{data}(x) + p_g(x)} \right]) dx = \\
 &\int_x (p_{data}(x)(-\log(2)) + p_g(x)(-\log(2))) dx + \\
 &\int_x (p_{data}(x) \log \left[\frac{p_{data}(x)}{\frac{p_{data}(x) + p_g(x)}{2}} \right] + p_g(x) \log \left[\frac{p_g(x)}{\frac{p_{data}(x) + p_g(x)}{2}} \right]) dx = \\
 &\int_x (p_{data}(x) + p_g(x))(-\log(2)) dx + \\
 &2 \left[\frac{1}{2} D_{KL}(p_r(x) \parallel \frac{p_{data}(x) + p_g(x)}{2}) + \frac{1}{2} D_{KL}(p_g(x) \parallel \frac{p_{data}(x) + p_g(x)}{2}) \right] \rightarrow \\
 \min_G \mathbf{V}(\mathbf{D}^*, \mathbf{G}) &= \mathbf{D}_{JS}(\mathbf{p}_r(\mathbf{x}) \parallel \mathbf{p}_g(\mathbf{x})) - \log(4)
 \end{aligned}$$

Assuming a fixed $G(z)$, and optimal $D^*(x) \rightarrow$

$$D^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_g(x)}$$

$$\min_G V(D^*, G) = D_{JS}(p_r(x) || p_g(x)) - \log(4)$$

So minimizing our GAN cost is the same as minimizing the Jensen-Shannon Distance between $p_r(x)$ and $p_{data}(x)$:

$$D_{JS}(p_r(x) || p_g(x))$$

For our GAN optimal point, our generator $G^*(z)$ matches the real data:

$$p_r(x) = p_{data}(x)$$

$$V(D^*, G^*) = -\log(4)$$

Training-Objective Function:

- WGAN/WGAN-GP.
- Relativistic GAN.
- Geometric GAN.
- f-GAN.
- Spectral Normalization GAN.
- BigGAN.

Structural Networks Changes:

- Self-Attention GAN.
- StyleGAN.
- ProGAN.
- DCGAN.

Image Quality and Diversity:

- ProGAN/StyleGAN.
- BigGAN.
- Self-Attention GAN.

Representation Learning:

- InfoGAN.
- BiGAN/BigBiGAN.
- StyleGAN/StyleGAN2.

Transfer Learning:

- Star-GAN.
- Cycle-GAN.

ImageNet: 1M images and 20K classes.



Figure 4: Generated and uncured images from a BigGAN. [2]

This person does not exist.



Figure 5: Generated and uncured images from a StyleGAN2. Karras et al. 2020.




Figure 6: Style mixing on StyleGAN. Karras et al. 2018

Representation learning on ImageNet (1M images and 20K classes)



Architecture	Top-1	Top-5
ResNet-50	76.3	93.1
ResNet-101	77.8	93.8
RevNet-50	71.8	90.5
RevNet-50 $\times 2$	74.9	92.2
RevNet-50 $\times 4$	76.6	93.1

Method	Architecture	Feature	Top-1	Top-5
BigBiGAN (ours)	ResNet-50	AvePool	55.4	77.4
	ResNet-50	BN+CReLU	56.6	78.6
	RevNet-50 $\times 4$	AvePool	60.8	81.4
	RevNet-50 $\times 4$	BN+CReLU	61.3	81.9

Figure 7: Image encoding/reconstruction and representation learning accuracy on BigBiGAN. Donahue et al. 2019. []

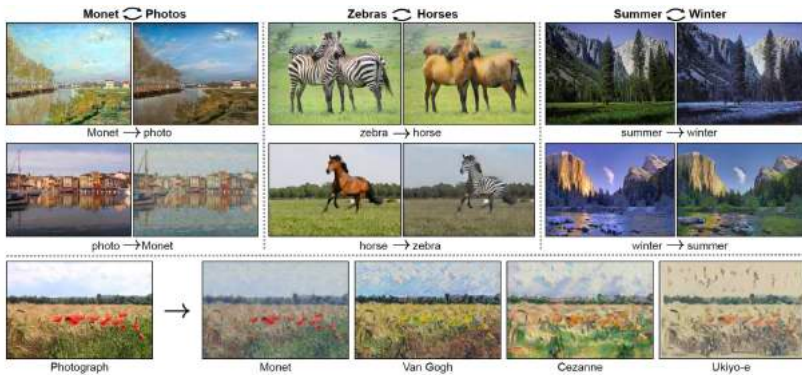


Figure 8: Cycle-GAN Zhu et al. 2017

- GANs:
 - Maximum Likelihood Models vs GANs.
- GAN Problems:
 - Non-convergence.
 - Diminished Gradient.
 - Unstable Gradients.
 - Mode Collapse.
- Evaluation:
 - Inception Score.
 - Frechet Inception Distance.
- General lines of research:
 - Stabilizing training and Improving learning.
 - Image Quality and Variety.
 - Representation learning.
 - Structural Changes.
 - Transfer Learning.

- **Further dip into GANs:** Myself - 2020
- **Another GAN introduction and models review:** Myself - 2020
- **Brief introduction to GANs:** Jonathan Hui - Medium
- **GAN models reading list, this is a good introduction to relevant models:** Jonathan Hui - Medium
- **Goodfellow NIPS 2016 Tutorial:**
 - Paper
 - Video
- **Play with Generative Adversarial Networks in your browser:** [Link](#)
- **Pytorch DCGAN Tutorial:** [Link](#)

- [1] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron C. Courville, and Yoshua Bengio.
Generative adversarial nets.
In Zoubin Ghahramani, Max Welling, Corinna Cortes, Neil D. Lawrence, and Kilian Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27: Annual Conference on Neural Information Processing Systems 2014, December 8-13 2014, Montreal, Quebec, Canada*, pages 2672–2680, 2014.
- [2] Andrew Brock, Jeff Donahue, and Karen Simonyan.
Large scale GAN training for high fidelity natural image synthesis.
CoRR, abs/1809.11096, 2018.